


Linear Probe Accuracy: Word vs Speaker Classification



Note: Untrained GAP word probe exceeds chance due to logistic regression exploiting spurious correlations in the compressed 128-dim random projection. Trained GAP word probe (91.1%) represents genuine learned word discrimination. Speaker probe shows suppression at Block 2 with partial recovery at GAP.